

In [19]:

```
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

%matplotlib inline

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("whitegrid")
sns.set_context("notebook")
#sns.set_context("poster")
```

Feature Selection

The element that has the biggest impact in the quality of your model is data features. You can only include in your model the attributes that you have and if they are not relevant, partially relevant or don't capture the causality relationships behind the model, or introduce other relationships that correspond to other causes different from the ones that you want to investigate, then you'll have a poor model.

Selecting the relevant features that add to your model is therefore of the utmost importance.

In this notebook we will deal with four approaches:

- 1) Univariate Selection.
- 2) Recursive feature elimination.
- 3) PCA - Principal Component Analysis.
- 4) Estimating feature importance.

Feature selection is a process where you select those features in your data that contribute most to the variable of interest. Irrelevant features decrease the accuracy of many models because you try to adjust on noise, this is particularly important in the case of linear models, such as linear and logistic regressions, where all features are always taken into account. Three are the main benefits of feature selection:

- 1) Reduces overfitting. Less redundant data implies less decisions made on noise.
- 2) Improves accuracy. Less misleading data results in a more accurate model.
- 3) Reduces training time. Less data implies faster training.

Scikitlearn has a nice and short article on feature selection where you can learn more https://scikit-learn.org/stable/modules/feature_selection.html (https://scikit-learn.org/stable/modules/feature_selection.html)

Again we will use the Pima Indians onset of diabetes dataset.



In this exercise we will use one of the traditional Machine Learning dataset, the Pima Indians diabetes dataset.

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger database. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

Content The datasets consists of several medical predictor variables and one target variable, **Outcome**. Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

- Pregnancies
- Glucose
- BloodPressure
- SkinThickness
- Insulin
- BMI
- DiabetesPedigreeFunction (scores de likelihood of diabetes based on family history)
- Age
- Outcome

In [20]:

```
# Load the Pima indians dataset and separate input and output components

from numpy import set_printoptions
set_printoptions(precision=3)

filename="pima-indians-diabetes.data.csv"
names=["pregnancies", "glucose", "pressure", "skin", "insulin", "bmi", "pedi", "age", "outcome"]
p_indians=pd.read_csv(filename, names=names)
p_indians.head()

# First we separate into input and output components
array=p_indians.values
X=array[:,0:8]
Y=array[:,8]
X
pd.DataFrame(X).head()
```

Out[20]:

	pregnancies	glucose	pressure	skin	insulin	bmi	pedi	age	outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Out[20]:

```
array([[ 6. , 148. , 72. , ..., 33.6 , 0.627, 50. ],
       [ 1. , 85. , 66. , ..., 26.6 , 0.351, 31. ],
       [ 8. , 183. , 64. , ..., 23.3 , 0.672, 32. ],
       ...,
       [ 5. , 121. , 72. , ..., 26.2 , 0.245, 30. ],
       [ 1. , 126. , 60. , ..., 30.1 , 0.349, 47. ],
       [ 1. , 93. , 70. , ..., 30.4 , 0.315, 23. ]])
```

Out[20]:

	0	1	2	3	4	5	6	7
0	6.0	148.0	72.0	35.0	0.0	33.6	0.627	50.0
1	1.0	85.0	66.0	29.0	0.0	26.6	0.351	31.0
2	8.0	183.0	64.0	0.0	0.0	23.3	0.672	32.0
3	1.0	89.0	66.0	23.0	94.0	28.1	0.167	21.0
4	0.0	137.0	40.0	35.0	168.0	43.1	2.288	33.0

Univariate Selection

One approach is to use statistical tests for example the Pearson Chi-Squared χ^2 is commonly used to select the most significant features.

We will use the **SelectKBest** class in scikit-learn.

In [21]:

```
# Univariate selection using Chi-squared
set_printoptions(precision=3)
p_indians.head()

from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2

# feature selection (we select the 4 best)
test = SelectKBest(score_func=chi2, k=4)
fit = test.fit(X,Y)
print("Scores")

fit.scores_

print("The 4 attributes with the highest scores are: glucose, insulin, bmi and age ")
print()

features=fit.transform(X)
features[0:5,:]
```

Out[21]:

	pregnancies	glucose	pressure	skin	insulin	bmi	pedi	age	outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Scores

Out[21]:

```
array([ 111.52 , 1411.887,   17.605,   53.108, 2175.565,  127.669,
        5.393,  181.304])
```

The 4 attributes with the highest scores are: glucose, insulin, bmi and age

Out[21]:

```
array([[148. ,   0. ,  33.6,  50. ],
       [ 85. ,   0. ,  26.6,  31. ],
       [183. ,   0. ,  23.3,  32. ],
       [ 89. ,  94. ,  28.1,  21. ],
       [137. , 168. ,  43.1,  33. ]])
```

In []:

Recursive Feature Elimination

This is a very intuitive approach. It consist on recursively removing attributes and building a model with those attributes remaining. It uses the model accuracy to identify which attributes or combination of attributes contribute the most.

We will use it with a logistic regression, but the choice of algorithm doesn't matter too much as long as your are consistent.

Recursive Feature Elimination uses the **RFE** class.

In [22]:

```
# Recursive Feature Elimiantion

from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression

p_indians.head()

#Logistic regression
model = LogisticRegression(solver='liblinear')

rfe = RFE(model, 3) # we want to find the 3 top features
fit = rfe.fit(X, Y)

print(f'Number of features {fit.n_features_:d}')
print(f'Selected features {fit.support_}')
print(f'Ranking of features {fit.ranking_}')
print()
print("Top features seem to be pregnancies, bmi, and pedi(Diabetes Pedigree Function)")
```

Out[22]:

	pregnancies	glucose	pressure	skin	insulin	bmi	pedi	age	outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Number of features 3

Selected features [True False False False False True True False]

Ranking of features [1 2 3 5 6 1 1 4]

Top features seem to be pregnancies, bmi, and pedi(Diabetes Pedigree Function)

In []:

Mission 1

For this and the next mission we will use data from Kaggle In concrete from the World University Rankings Competition <https://www.kaggle.com/mylesoneill/world-university-rankings> (<https://www.kaggle.com/mylesoneill/world-university-rankings>)

a) Using the Shanghai rankings find the top 3 most important features to explain them with both univariate and recursive (in recursive because we are using log regression create an output variable of being in the top 50 or not).

b) Same for the Times ranking.

c) Does it change if we choose the top 10 or top 100?

In [23]:

```
# Clean shanghai dataset function

def shangai_clean(x):
    # Read excel file and sort by total score
    shangai = pd.read_excel("shanghaiData.xlsx").sort_values(by = "total_score", ascending

    # Filter by the latest year
    shangai = shangai[shangai["year"] == shangai["year"].max()]

    # Simplify dataframe with only explanatory variables and drop null values
    shangai.drop(["world_rank", "university_name", "national_rank", "year", "total_score"],
    shangai.dropna(inplace = True)

    # Code the top 50 universities
    array_ref = (np.arange(len(shangai)) < x)
    shangai["top_50"] = array_ref
    code = {True:1.0, False:0.0}
    shangai["top_50"] = shangai["top_50"].map(code)

    # Return the array
    return shangai
```

In [24]:

```
# Automating column printing function for recursive feature elimination

def col_list(ranking, column_names):
    var_s = ranking
    top_cols = []
    counter = 0
    for x in var_s:
        if var_s[counter] == 1:
            top_cols.append(column_names[counter])
            counter +=1
    return top_cols
```

In [25]:

```
# Automatating column printing function for univariate analysis

def col_dic(scores):
    import math
    indices = np.arange(len(scores))
    first_in = second_in = third_in = 0
    first = second = third = -math.inf
    for x in indices:
        if scores[x] > first:
            second = first
            second_in = first_in
            first = scores[x]
            first_in = x
        elif scores[x] > second:
            third = second
            third_in = second_in
            second = scores[x]
            second_in = x
        elif scores[x] > third:
            third = scores[x]
            third_in = x
    return {"first":first_in, "second":second_in, "third": third_in}
```

In [32]:

```

# Shangai top-50 analysis

shangai_treat = shangai_clean(50)
shan_cols = shangai_treat.columns

array = shangai_treat.values

shan_x = array[:,0:6]
shan_y = array[:,6]

# Run Logistic Regression

model = LogisticRegression(solver='liblinear')
rfe = RFE(model, 3)# we want to find the 3 top features
rec = rfe.fit(shan_x, shan_y)

# Automate column outputs for RFE

top_cols = col_list(rec.ranking_, shan_cols)

print('\033[1m' + 'Recursive Feature Elimination' '\033[0m')
print()
print(f'Number of features {rec.n_features_:d}')
print(f'Selected features {rec.support_}')
print(f'Ranking of features {rec.ranking_}')
print()
print(f"Top features are: {top_cols[0]}, {top_cols[1]}, {top_cols[2]}")

# Run univariate analysis

test = SelectKBest(score_func=chi2, k=3)
uni = test.fit(shan_x, shan_y)

print()
print('\033[1m' + 'Univariate Analysis' '\033[0m')
print()
uni.scores_

# Automate column outputs for univariate analysis

cols = col_dic(uni.scores_)
print("The top features are: " + shan_cols[cols["first"]] + ", " + shan_cols[cols["second"]])

features=uni.transform(shan_x)
features[0:5,:]

```

Recursive Feature Elimination

Number of features 3

Selected features [False True False True False True]

Ranking of features [4 1 2 1 3 1]

Top features are: award, ns, pcp

Univariate Analysis

Out[32]:




```
array([3533.277, 8078.695, 2861.682, 2551.689, 599.288, 528.423])
```

The top features are: award, alumni, hici

Out[32]:

```
array([[100. , 100. , 100. ],
       [ 40.7,  89.6,  80.1],
       [ 68.2,  80.7,  60.6],
       [ 65.1,  79.4,  66.1],
       [ 77.1,  96.6,  50.8]])
```

In [27]:

```
def times_clean(a):

    # Read the csv file and sort by total score ind descending order
    times = pd.read_csv("timesData.csv").sort_values(by = "total_score", ascending = False)

    # Filter by the latest year
    times = times[times["year"] == 2016]

    # Simplify the table by dropping non explanatory variables
    times.drop(["world_rank", "university_name", "country", "year", "total_score"], axis =

    # Drop null values
    times.dropna(inplace = True)

    # Convert all other columns to float type
    times["international"] = pd.to_numeric(times["international"], errors = "coerce")
    times["income"] = pd.to_numeric(times["income"], errors = "coerce")
    times["female_male_ratio"] = pd.to_numeric(times["female_male_ratio"].apply(lambda d: c
    times["international_students"] = pd.to_numeric(times["international_students"].apply(I

    student_list = []

    for x in range(len(times)):
        student_value = times["num_students"].iloc[x].replace(",",".")
        student_list.append(student_value)

    times["num_students"] = pd.to_numeric(student_list, errors = "coerce")

    # Add a new column with the actual ranking
    array_ref = (np.arange(len(times)) < a)
    times["top_50"] = array_ref
    code = {True:1.0, False:0.0}
    times["top_50"] = times["top_50"].map(code)

    times.dropna(inplace = True)

    return times
```

In [28]:

```

# Times top-50 analysis

times_treat = times_clean(50)
times_cols = times_treat.columns
array = times_treat.values

times_x = array[:,0:9]
times_y = array[:,9]

# Run Logistic Regression

model = LogisticRegression(solver='liblinear')
rfe = RFE(model, 3)
rec = rfe.fit(times_x, times_y)

# Automate column outputs for RFE

top_cols = col_list(rec.ranking_, times_cols)

print('\033[1m' + 'Recursive Feature Elimination' '\033[0m')
print()
print(f'Number of features {rec.n_features_:d}')
print(f'Selected features {rec.support_}')
print(f'Ranking of features {rec.ranking_}')
print()
print(f"Top features are {top_cols[0]}, {top_cols[1]}, and {top_cols[2]}")

# Run univariate analysis

test = SelectKBest(score_func=chi2, k=3)
uni = test.fit(times_x, times_y)

print()
print('\033[1m' + 'Univariate Analysis' '\033[0m')
print()
uni.scores_

cols = col_dic(uni.scores_)
print("The top features are: " + times_cols[cols["first"]] + ", " + times_cols[cols["second"]])

features=uni.transform(times_x)
features[0:9,:]

```

Recursive Feature Elimination

Number of features 3

Selected features [False False True False False False True False True]

Ranking of features [2 6 1 7 3 5 1 4 1]

Top features are research, student_staff_ratio, and female_male_ratio

Univariate Analysis

Out[28]:

array([2.552e+03, 5.348e+02, 4.891e+03, 1.552e+03, 3.463e+02, 9.933e-02,

8.197e+01, 5.372e+02, 8.368e-01])



The top features are: research, teaching, citations

Out[28]:

```
array([[95.6, 97.6, 99.8],
       [86.5, 98.9, 98.8],
       [92.5, 96.2, 99.9],
       [88.2, 96.7, 97. ],
       [89.4, 88.6, 99.7],
       [85.1, 91.9, 99.3],
       [83.3, 88.5, 96.7],
       [77. , 95. , 91.1],
       [85.7, 88.9, 99.2]])
```

In [29]:

```
# c) Does it change if we choose the top 10 or top 100?
```

In [30]:

```

# Shangai top-10 analysis

shangai_10 = shangai_clean(10)
shan_cols = shangai_treat.columns
array_10 = shangai_10.values

shan_x10 = array_10[:,0:6]
shan_y10 = array_10[:,6]

# Run Logistic Regression

model = LogisticRegression(solver='liblinear')
rfe = RFE(model, 3)# we want to find the 3 top features
rec = rfe.fit(shan_x10, shan_y10)

# Automate column outputs for RFE

top_cols_10 = col_list(rec.ranking_, shan_cols)

print('\033[1m' + 'Recursive Feature Elimination' '\033[0m')
print()
print(f'Number of features {rec.n_features_:d}')
print(f'Selected features {rec.support_}')
print(f'Ranking of features {rec.ranking_}')
print()
print(f"Top features are: {top_cols_10[0]}, {top_cols_10[1]}, {top_cols_10[2]}")

# Run univariate analysis

test = SelectKBest(score_func=chi2, k=3)
uni = test.fit(shan_x10, shan_y10)

print()
print('\033[1m' + 'Univariate Analysis' '\033[0m')
print()
uni.scores_

# Automate column outputs for univariate analysis

cols_10 = col_dic(uni.scores_)
print("The top features are: " + shan_cols[cols_10["first"]] + ", " + shan_cols[cols_10["se

features=uni.transform(shan_x10)
features[0:5,:]
```

Recursive Feature Elimination

Number of features 3

Selected features [True False True False True False]

Ranking of features [1 4 1 3 1 2]

Top features are: alumni, hici, pub

Univariate Analysis

Out[30]:

```
array([3839.867, 7499.15 , 1446.554, 1397.945, 168.597, 681.022])
```



The top features are: award, alumni, hici

Out[30]:

```
array([[100. , 100. , 100. ],
       [ 40.7,  89.6,  80.1],
       [ 68.2,  80.7,  60.6],
       [ 65.1,  79.4,  66.1],
       [ 77.1,  96.6,  50.8]])
```

In [9]:

```

# Shangai top 100 analysis

shangai_100 = shangai_clean(100)
shan_cols = shangai_treat.columns
array_100 = shangai_100.values

shan_x100 = array_100[:,0:6]
shan_y100 = array_100[:,6]

# Run Logistic Regression

model = LogisticRegression(solver='liblinear')
rfe = RFE(model, 3)# we want to find the 3 top features
rec = rfe.fit(shan_x100, shan_y100)

# Automate column outputs for RFE

top_cols = col_list(rec.ranking_, shan_cols)

print('\033[1m' + 'Recursive Feature Elimination' '\033[0m')
print()
print(f'Number of features {rec.n_features_:d}')
print(f'Selected features {rec.support_}')
print(f'Ranking of features {rec.ranking_}')
print()
print(f"Top features are: {top_cols[0]}, {top_cols[1]}, {top_cols[2]}")

# Run univariate analysis

test = SelectKBest(score_func=chi2, k=3)
uni = test.fit(shan_x100, shan_y100)

print()
print('\033[1m' + 'Univariate Analysis' '\033[0m')
print()
uni.scores_

# Automate column outputs for univariate analysis

cols = col_dic(uni.scores_)
print("The top features are: " + shan_cols[cols["first"]] + ", " + shan_cols[cols["second"]])

features=uni.transform(shan_x100)
features[0:5,:]
```

Recursive Feature Elimination

Number of features 3

Selected features [True True True False False False]

Ranking of features [1 1 1 3 4 2]

Top features are: alumni, award, hici

Univariate Analysis

Out[9]:

array([3589.923, 7203.251, 2934.493, 2328.666, 665.928, 577.754])

The top features are: award, alumni, hici

Out[9]:

```
array([[100. , 100. , 100. ],
       [ 40.7,  89.6,  80.1],
       [ 68.2,  80.7,  60.6],
       [ 65.1,  79.4,  66.1],
       [ 77.1,  96.6,  50.8]])
```

In [14]:

```

# Times top-10 analysis

times_treat = times_clean(10)
times_cols = times_treat.columns
array_10 = times_treat.values

times_x10 = array_10[:,0:9]
times_y10 = array_10[:,9]

# Run Logistic Regression

model = LogisticRegression(solver='liblinear')
rfe = RFE(model, 3)
rec = rfe.fit(times_x10, times_y10)

# Automate column outputs for RFE

top_cols = col_list(rec.ranking_, times_cols)

print('\033[1m' + 'Recursive Feature Elimination' '\033[0m')
print()
print(f'Number of features {rec.n_features_:d}')
print(f'Selected features {rec.support_}')
print(f'Ranking of features {rec.ranking_}')
print()
print(f"Top features are {top_cols[0]}, {top_cols[1]}, and {top_cols[2]}")

# Run univariate analysis

test = SelectKBest(score_func=chi2, k=3)
uni = test.fit(times_x10, times_y10)

print()
print('\033[1m' + 'Univariate Analysis' '\033[0m')
print()
uni.scores_

cols = col_dic(uni.scores_)
print("The top features are: " + times_cols[cols["first"]] + ", " + times_cols[cols["second"]])

features=uni.transform(times_x10)
features[0:9,:]
```

Recursive Feature Elimination

Number of features 3

Selected features [False False True False False False True False True]

Ranking of features [2 5 1 4 6 3 1 7 1]

Top features are research, student_staff_ratio, and female_male_ratio

Univariate Analysis

Out[14]:

```
array([ 966.246,  217.918, 1522.134,  410.647,  130.602,   55.884,
        53.895,  247.092,   14.761])
```


The top features are: research, teaching, citations

Out[14]:

```
array([[95.6, 97.6, 99.8],  
       [86.5, 98.9, 98.8],  
       [92.5, 96.2, 99.9],  
       [88.2, 96.7, 97. ],  
       [89.4, 88.6, 99.7],  
       [85.1, 91.9, 99.3],  
       [83.3, 88.5, 96.7],  
       [77. , 95. , 91.1],  
       [85.7, 88.9, 99.2]])
```

In [15]:

```

# Times top-100 analysis

times_treat = times_clean(100)
times_cols = times_treat.columns
array_100 = times_treat.values

times_x100 = array_100[:,0:9]
times_y100 = array_100[:,9]

# Run Logistic Regression

model = LogisticRegression(solver='liblinear')
rfe = RFE(model, 3)
rec = rfe.fit(times_x100, times_y100)

# Automate column outputs for RFE

top_cols = col_list(rec.ranking_, times_cols)

print('\033[1m' + 'Recursive Feature Elimination' '\033[0m')
print()
print(f'Number of features {rec.n_features_:d}')
print(f'Selected features {rec.support_}')
print(f'Ranking of features {rec.ranking_}')
print()
print(f"Top features are {top_cols[0]}, {top_cols[1]}, and {top_cols[2]}")

# Run univariate analysis

test = SelectKBest(score_func=chi2, k=3)
uni = test.fit(times_x100, times_y100)

print()
print('\033[1m' + 'Univariate Analysis' '\033[0m')
print()
uni.scores_

cols = col_dic(uni.scores_)
print("The top features are: " + times_cols[cols["first"]] + ", " + times_cols[cols["second"]])

features=uni.transform(times_x100)
features[0:9,:]

```

Recursive Feature Elimination

Number of features 3

Selected features [False False True True False False False False True]

Ranking of features [5 6 1 1 3 7 4 2 1]

Top features are research, citations, and female_male_ratio

Univariate Analysis

Out[15]:

```

array([2.669e+03, 9.792e+02, 5.515e+03, 2.756e+03, 4.808e+02, 5.439e+01,
       6.599e+01, 7.550e+02, 4.563e-01])

```

The top features are: research, citations, teaching

Out[15]:

```
array([[95.6, 97.6, 99.8],
       [86.5, 98.9, 98.8],
       [92.5, 96.2, 99.9],
       [88.2, 96.7, 97. ],
       [89.4, 88.6, 99.7],
       [85.1, 91.9, 99.3],
       [83.3, 88.5, 96.7],
       [77. , 95. , 91.1],
       [85.7, 88.9, 99.2]])
```

In [15]:

```
print("As we can see for both datasets, scores and more important features have changed for
print()
print("This is due to the fact that the output variable is different for top-10 and top-100
print()
print("The factors determining whether a university is elite (top-10) can be different from
print("university is very good (top-100).")
print()
print("Hence, different variables should explain the variance in the different output varia
```

As we can see for both datasets, scores and more important features have changed for the top-10 or top-100.

This is due to the fact that the output variable is different for top-10 and top-100.

The factors determining whether a university is elite (top-10) can be different from the factors determining whether a university is very good (top-100).

Hence, different variables should explain the variance in the different output variable.

Principal Component Analysis

Principal Component Analysis is a data reduction technique using linear algebra. The idea here is to "compress" several dimensions into principal components.

One problem of PCA is the explainability. Once you compressed the attributes into principal components you can no longer refer to them individually establishing causality links or relationships.

A property of PCA is that you can choose the number of dimensions or principal components. In our example we will select 3 principal components.

For Principal Component Analysis you use the **PCA** class.

In [16]:

```

from sklearn.decomposition import PCA

p_indians.head()

#PCA
pca = PCA(n_components=3)
pca_fit = pca.fit(X)

print(f"Explained variance: {pca_fit.explained_variance_ratio_}")
print()

np.set_printoptions(formatter={'float': '{: 0.3f}'.format})
print("Principal Components have little resemblance to the source data attributes")
print()
print(pca_fit.components_)

```

Out[16]:

	pregnancies	glucose	pressure	skin	insulin	bmi	pedi	age	outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Explained variance: [0.889 0.062 0.026]

Principal Components have little resemblance to the source data attributes

```

[[-0.002  0.098  0.016  0.061  0.993  0.014  0.001 -0.004]
 [-0.023 -0.972 -0.142  0.058  0.095 -0.047 -0.001 -0.140]
 [-0.022  0.143 -0.922 -0.307  0.021 -0.132 -0.001 -0.125]]

```

In []:

Feature Importance

One of the added features of tree based algorithms is that they can be used to estimate the importance of each feature and use it to refine the model to different levels depending on where we want to situate ourselves in the tension between explainability and accuracy.

In this example we are going to use the `ExtraTreesClassifier`, but the technique is commonly used in all tree algorithms.

For this example of assessing feature importance with trees we will use the **ExtraTreesClassifier** class.

In [17]:

```

from sklearn.ensemble import ExtraTreesClassifier

p_indians.head()

model = ExtraTreesClassifier(n_estimators=100)
model.fit(X,Y)

print(model.feature_importances_)

```

Out[17]:

	pregnancies	glucose	pressure	skin	insulin	bmi	pedi	age	outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Out[17]:

```

ExtraTreesClassifier(bootstrap=False, class_weight=None, criterion='gini',
    max_depth=None, max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
    oob_score=False, random_state=None, verbose=0, warm_start=False)

[ 0.109  0.231  0.098  0.078  0.076  0.146  0.119  0.143]

```

In []:

Mission 2

- Using the Shangai Data find the top attributes with a tree classifier for top-10, top-50 and top-100.
- Same for the Times ranking.

In [18]:

```
# a) Using the Shanghai Data find the top attributes with a tree classifier for top-10, top-50, top-100

array_100 = shangai_clean(100).values
array_50 = shangai_clean(50).values
array_10 = shangai_clean(10).values

x_100 = array_100[:,0:6]
y_100 = array_100[:,6]

x_50 = array_50[:,0:6]
y_50 = array_50[:,6]

x_10 = array_10[:,0:6]
y_10 = array_10[:,6]

# Model with top 100 estimators

print('\033[1m' + 'Shanghai dataset with top-100 universities' '\033[0m')

model_100 = ExtraTreesClassifier(n_estimators=100)
model_100.fit(x_100,y_100)

print(model_100.feature_importances_)
indices_100 = model_100.feature_importances_.argsort()[::-1][:3]
print()
print("The top 3 features are: " + shan_cols[indices_100[0]] + ", " + shan_cols[indices_100[1]] + ", " + shan_cols[indices_100[2]])
print()

# Model with top 50 estimators

print('\033[1m' + 'Shanghai dataset with top-50 universities' '\033[0m')

model_50 = ExtraTreesClassifier(n_estimators=100)
model_50.fit(x_50,y_50)

print(model_50.feature_importances_)
indices_50 = model_50.feature_importances_.argsort()[::-1][:3]
print()
print("The top 3 features are: " + shan_cols[indices_50[0]] + ", " + shan_cols[indices_50[1]] + ", " + shan_cols[indices_50[2]])
print()

# Model with top 10 estimators

print('\033[1m' + 'Shanghai dataset with top-10 universities' '\033[0m')

model_10 = ExtraTreesClassifier(n_estimators=100)
model_10.fit(x_10,y_10)

print(model_10.feature_importances_)
indices_10 = model_10.feature_importances_.argsort()[::-1][:3]
print()
print("The top 3 features are: " + shan_cols[indices_10[0]] + ", " + shan_cols[indices_10[1]] + ", " + shan_cols[indices_10[2]])
print()
```

Shanghai dataset with top-100 universities

Out[18]:

```
ExtraTreesClassifier(bootstrap=False, class_weight=None, criterion='gini',
                    max_depth=None, max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                    oob_score=False, random_state=None, verbose=0, warm_start=False)
```

```
[ 0.099  0.239  0.265  0.177  0.114  0.105]
```

The top 3 features are: alumni, pcg, pub

Shangai dataset with top-50 universities

Out[18]:

```
ExtraTreesClassifier(bootstrap=False, class_weight=None, criterion='gini',
                    max_depth=None, max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                    oob_score=False, random_state=None, verbose=0, warm_start=False)
```

```
[ 0.088  0.217  0.243  0.254  0.110  0.087]
```

The top 3 features are: pcg, alumni, pub

Shangai dataset with top-10 universities

Out[18]:

```
ExtraTreesClassifier(bootstrap=False, class_weight=None, criterion='gini',
                    max_depth=None, max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                    oob_score=False, random_state=None, verbose=0, warm_start=False)
```

```
[ 0.099  0.239  0.265  0.177  0.114  0.105]
```

The top 3 features are: alumni, pcg, pub

In [23]:

```

# b) Same for the Times ranking.

array_100 = times_clean(100).values
array_50 = times_clean(50).values
array_10 = times_clean(10).values

x_100 = array_100[:,0:9]
y_100 = array_100[:,9]

x_50 = array_50[:,0:9]
y_50 = array_50[:,9]

x_10 = array_10[:,0:9]
y_10 = array_10[:,9]

# Model with top 100 estimators

print('\033[1m' + 'Times dataset with top-100 universities' '\033[0m')

model_100 = ExtraTreesClassifier(n_estimators=100)
model_100.fit(x_100,y_100)

indices_100 = model_100.feature_importances_.argsort()[::-1][:3]
print("The top 3 features are: " + times_cols[indices_100[0]] + ", " + times_cols[indices_100[1]] + ", " + times_cols[indices_100[2]])
print()

print(model_100.feature_importances_)

# Model with top 50 estimators

print()
print('\033[1m' + 'Times dataset with top-50 universities' '\033[0m')

model_50 = ExtraTreesClassifier(n_estimators=100)
model_50.fit(x_50,y_50)

indices_50 = model_50.feature_importances_.argsort()[::-1][:3]
print("The top 3 features are: " + times_cols[indices_50[0]] + ", " + times_cols[indices_50[1]] + ", " + times_cols[indices_50[2]])
print()

print(model_50.feature_importances_)

# Model with top 10 estimators

print()
print('\033[1m' + 'Times dataset with top-10 universities' '\033[0m')

model_10 = ExtraTreesClassifier(n_estimators=100)
model_10.fit(x_10,y_10)

indices_10 = model_10.feature_importances_.argsort()[::-1][:3]
print("The top 3 features are: " + times_cols[indices_10[0]] + ", " + times_cols[indices_10[1]] + ", " + times_cols[indices_10[2]])
print()

print(model_10.feature_importances_)

```

Times dataset with top-100 universities

Out[23]:

```
ExtraTreesClassifier(bootstrap=False, class_weight=None, criterion='gini',
                    max_depth=None, max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                    oob_score=False, random_state=None, verbose=0, warm_start=False)
```

The top 3 features are: female_male_ratio, student_staff_ratio, num_students

```
[ 0.258  0.049  0.368  0.171  0.037  0.024  0.022  0.047  0.022]
```

Times dataset with top-50 universities

Out[23]:

```
ExtraTreesClassifier(bootstrap=False, class_weight=None, criterion='gini',
                    max_depth=None, max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                    oob_score=False, random_state=None, verbose=0, warm_start=False)
```

The top 3 features are: num_students, student_staff_ratio, female_male_ratio

```
[ 0.368  0.035  0.390  0.090  0.029  0.017  0.018  0.036  0.018]
```

Times dataset with top-10 universities

Out[23]:

```
ExtraTreesClassifier(bootstrap=False, class_weight=None, criterion='gini',
                    max_depth=None, max_features='auto', max_leaf_nodes=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                    oob_score=False, random_state=None, verbose=0, warm_start=False)
```

The top 3 features are: income, student_staff_ratio, num_students

```
[ 0.258  0.049  0.368  0.171  0.037  0.024  0.022  0.047  0.022]
```

In []: